



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

3253 Analytic Techniques and Machine Learning

**Module 2: End to End Machine
Learning Project**



Course Plan

Module Titles

Module 1 – Introduction to Machine Learning

Module 2 – Current Focus: End to End Machine Learning Project

Module 3 – Classification

Module 4 – Clustering and Unsupervised Learning

Module 5 – Training Models and Feature Selection

Module 6 – Support Vector Machines

Module 7 – Decision Trees and Ensemble Learning

Module 8 – Dimensionality Reduction

Module 9 – Introduction to TensorFlow

Module 10 – Introduction to Deep Learning and Deep Neural Networks

Module 11 – Distributing TensorFlow, CNNs and RNNs

Module 12 – Final Assignment and Presentations (no content)



Learning Outcomes for this Module

- Perform all of the steps of building a simple machine learning project:
 - Conduct exploratory analysis and visualization
 - Prepare data
 - Set aside a testing set
 - Select and train a model
 - Deploy, monitor and maintain it



Topics for this Module

- **2.1** Loss functions
- **2.2** Feature scaling
- **2.3** Test sets
- **2.4** Geographical data
- **2.5** Correlation
- **2.6** Encoding features
- **2.7** Cross-validation
- **2.8** Searching for hyperparameters
- **2.9** Running the application
- **2.10** Resources and Wrap-up



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 1

Loss Functions

Notations

- m is the number of instances in dataset

- $x^{(1)} = \begin{pmatrix} -118.29 \\ 33.91 \\ 1416 \\ 38,372 \end{pmatrix}$

- \mathbf{X} is a matrix containing all feature values

- $\mathbf{X} = \begin{pmatrix} (x^1)^T \\ (x^{(2)})^T \\ \dots \\ (x^{100})^T \end{pmatrix} = \begin{pmatrix} -118.29 & 33.91 & 1416 & 38,372 \\ \vdots & \vdots & \vdots & \vdots \end{pmatrix}$

Predictions

- h is system's prediction function, aka a *hypothesis*
- When system is given an input $\mathbf{x}^{(i)}$, it outputs a predicted value $y_{\text{hat}}(i) = h(\mathbf{x}^{(i)})$
- For a regression task, will often use Root Mean Squared Error (RMS-E) or Mean Absolute Error (MAE) as the cost function

Performance Measures

- Root Mean Square Error (RMSE) measures standard deviation c of errors in prediction
- $$RMSE(X, h) = \sqrt{\frac{1}{m} \sum_{i=1}^m (h(x^{(i)}) - y^{(i)})^2}$$
- In housing example, assuming a normal distribution, $\sigma=50,000$ implies that
 - ~68% of system's predictions fall within \$50,000 of actual value
 - ~95% of predictions fall within \$100,000 of actual value

Performance Measures (cont'd)

- Mean Absolute Error (MAE) measures the average magnitude of the errors in a set of predictions, without considering their direction.
- It's the average over the test sample of the absolute differences between prediction and actual observation where all individual differences have equal weight.

$$\text{MAE} = \frac{1}{n} \sum_{j=1}^n |y_j - \hat{y}_j|$$



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 2

Feature Scaling

Feature Scaling

- ML algorithms usually perform poorly when input features have different scales
- Min-max scaling, aka *normalization*: values are shifted and rescaled so result is in range $[0,1]$
- *Standardization*:
 - Subtract mean (giving zero mean)
 - Divide by variance (unit mean)

Feature Scaling (cont'd)

- Standardization does not bound values to a specific range
- Can be problematic for certain ML methods — neural nets work well with input range [0,1]
- Much less affected by outliers than min-max
 - one large value will “crush” all other values into a very limited range



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 3

Test Sets

Create a Test Set

- *Data snooping bias* -- studying the test set might result in observing some interesting pattern, which when accounted for in model means that model will not generalize as well
- Pick some instances randomly from dataset randomly and set them aside
- Typically 20% of the dataset
- Important to set aside subset before training starts, don't randomly generate each run
 - over time, your algorithm will see whole dataset, which will invalidate the whole validation process

Create a Test Set (cont'd)

- Pick some instances randomly from dataset randomly and set them aside
- Typically 20% of the dataset
- Important to set aside subset before training starts, don't randomly generate each run
 - over time, your algorithm will see whole dataset, which will invalidate the whole validation process

Create a Test Set (cont'd)

- Using random sampling is usually acceptable if dataset is sufficiently large with respect to number of attributes
- For small datasets, purely random sampling may not work
- Example:
 - Population of USA composed of 51.3% females, 48.7% males
 - ensuring that a survey of the population maintains this ratio is referred to as *stratified sampling*

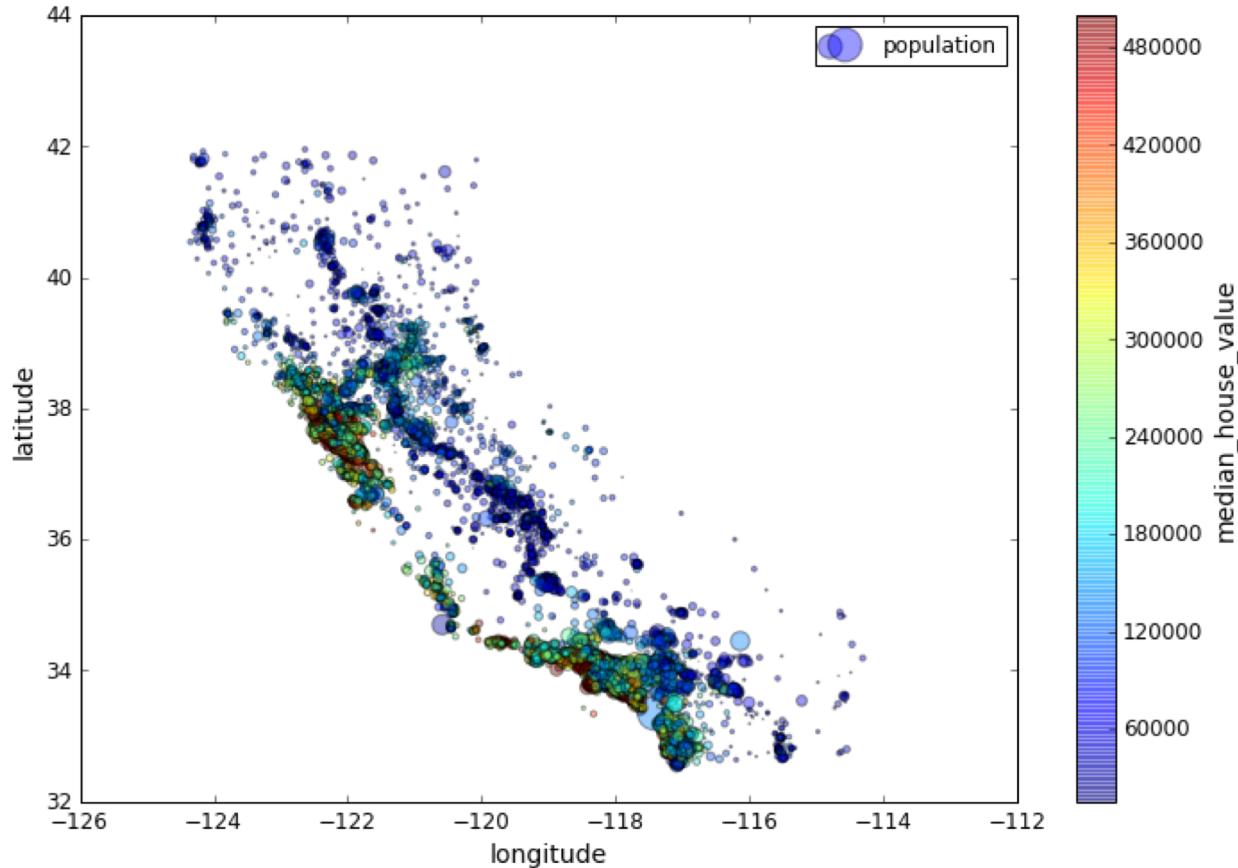


UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 4

Geographical Data

Visualizing Geographical Data



Visualizing Geographical Data (cont'd)

- As seen in previous image, strong correlation between population size and income
- Also related to distance to the coast
- Plotting data can be extremely useful to gain insights to test
- Scatterplots (as seen previously) can be helpful, histograms too
- Simple plotting may expose a relationship between variables



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

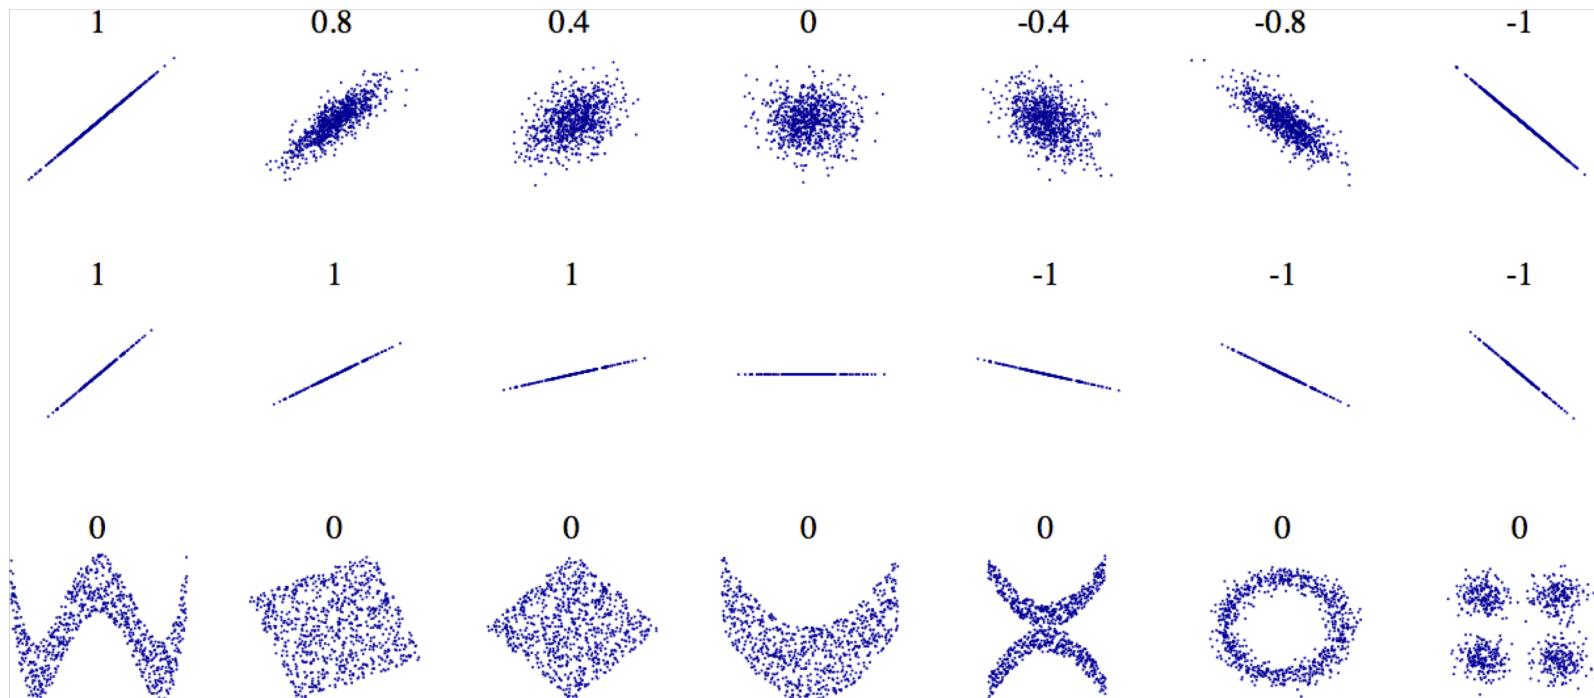
Module 2 – Section 5

Correlation

Correlation Coefficient

- Correlation is measure of the degree to which two variables have a linear relationship between them
- Correlation coefficient varies between [-1.0,1.0]
- Note that correlation coefficient only measures *linear* relationship between variables
- Pearson's correlation coefficient $\rho_{x,y} = \frac{cov(x,y)}{\sigma_x\sigma_y}$

Correlation Coefficient (cont'd)





Module 2 – Section 6

Encoding Features

One-Hot-Encoding

Sample	Category	Numerical	Sample	Human	Penguin	Octopus	Alien
1	Human	1	1	1	0	0	0
2	Human	1	2	1	0	0	0
3	Penguin	2	3	0	1	0	0
4	Octopus	3	4	0	0	1	0
5	Alien	4	5	0	0	0	1
6	Octopus	3	6	0	0	1	0
7	Alien	4	7	0	0	0	1

<https://www.quora.com/What-is-one-hot-encoding-and-when-is-it-used-in-data-science>



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 7

Cross-Validation

Cross Validation

- *K-fold cross validation* — randomly split into k distinct subsets called *folds*
- Trains and evaluates model k times, picking different fold each time for validation and training on the k-1 other folds
- Allows for estimate of model performance, but also estimate of how precise this estimate is (standard deviation)



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 8

Searching for Hyperparameters

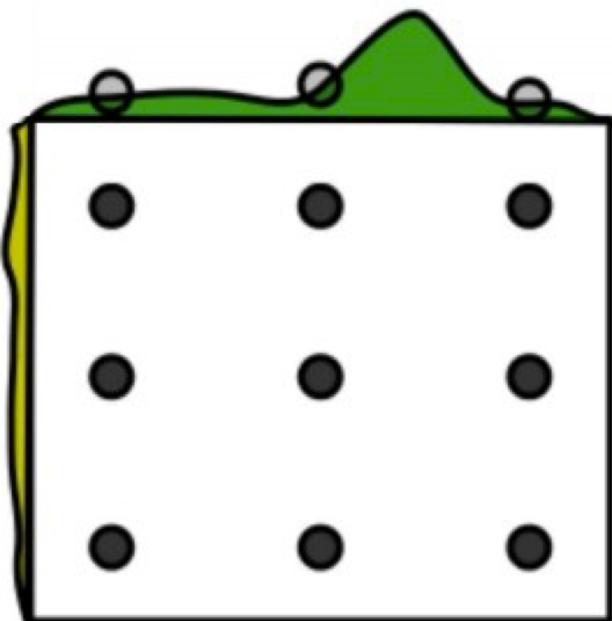
Searching for Hyperparameters

- Once a basic model has been created, need to search for optimal hyperparameters (learning rate, weights to balance loss terms, etc)
- *Randomized search*: useful for large search spaces
- Evaluates a given number of random combinations by selecting a random value for each hyperparam at each iteration
- Grid search: try every combination of the sets of hyperparameters provided

Grid Search

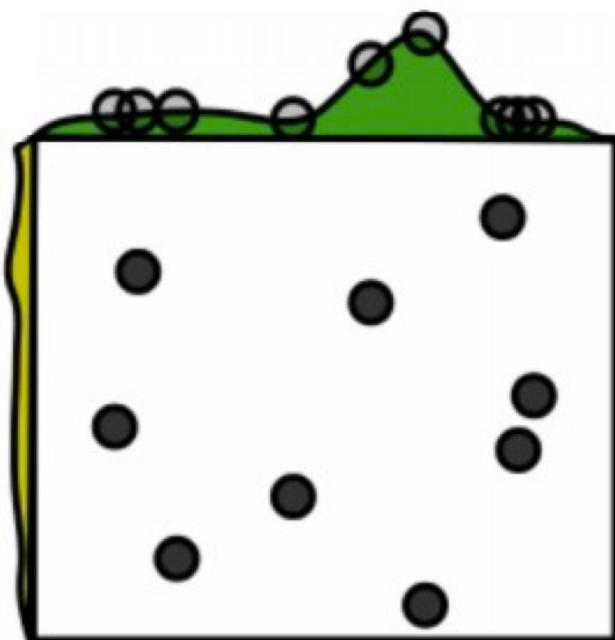
Grid Layout

Unimportant parameter



Random Layout

Unimportant parameter



Grid Search (cont'd)

- Each dict has keys that correspond to args for the model
- Associated dict is an exhaustive list of params to be tested
- Grid search finds best set of params

```
param_grid = [  
    {'n_estimators':[3, 10, 30],  
     'max_features':[2, 4, 6, 8]},  
    {'bootstrap':[False],  
     'n_estimators': [3, 10],  
     'max_features':[2, 3, 4]}]  
  
forest_reg =  
RandomForestRegressor()  
  
grid_search =  
GridSearchCV(forest_reg,  
            param_grid,  
            cv=5,  
            scoring='neg_mean_squared_error')
```



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 9

Running the Application

Launch, Monitor and Maintain System

- Need to write code to monitor performance
- Not just monitoring for breakage, but ability of model to make accurate predictions as data evolves (“data rot”)
- Make sure to monitor quality of inputted data: performance of model may degrade slowly
- Catching issues like malfunctioning sensor input or stale data from another team will allow for faster resolution of issues



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Module 2 – Section 10

Resources and Wrap-up

Homework

- Complete the notebook in the assignments section for this week
- Make sure you rename your notebook to
 - W2_UTORid.ipynb
 - Example: W2_adfasd01.ipynb

Next Class

- Classification
- Reading: Chapter 3

Attendance

- Please fill in the form
 - <https://goo.gl/forms/m0EVxGlkxB27LlsS2>

Follow us on social

Join the conversation with us online:

 [facebook.com/uoftscs](https://www.facebook.com/uoftscs)

 [@uoftscs](https://twitter.com/uoftscs)

 [linkedin.com/company/university-of-toronto-school-of-continuing-studies](https://www.linkedin.com/company/university-of-toronto-school-of-continuing-studies)

 [@uoftscs](https://www.instagram.com/uoftscs)



UNIVERSITY OF TORONTO
SCHOOL OF CONTINUING STUDIES

Any questions?



Thank You

Thank you for choosing the University of Toronto
School of Continuing Studies